

```
In [1]: import numpy as np
import pandas as pd

Import visualization libraries and set %matplotlib inline.

In [2]: import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
%matplotlib inline

Read in the csv file as a dataframe called df

In [3]: df = pd.read_csv('911.csv')

Check the info() of the df

In [4]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 99492 entries, 0 to 99491
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  --
0   lat                   99492 non-null  float64
1   lng                   99492 non-null  float64
2   desc                  99492 non-null  object
3   zip                   86637 non-null  float64
4   title                 99492 non-null  object
5   timeStamp             99492 non-null  object
6   twp                   99492 non-null  object
7   addr                  98773 non-null  object
8   e                     99492 non-null  int64
dtypes: float64(3), int64(1), object(5)
memory usage: 6.8+ MB

Check the head of df

In [5]: df.head(3)

Out [5]:
```

	lat	lng	desc	zip	title	timeStamp	twp	addr	e
0	40.297876	-75.561294	REINDEER CT & DEAD END- NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	2015-12-10 17:40:00	NEW HANOVER	REINDEER CT & DEAD END	1
1	40.250601	-75.264680	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	2015-12-10 17:40:00	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1
2	40.121182	-75.351375	HAWS AVE, NORRISTOWN, 2015-12-10 @ 14:39:21 SL...	19401.0	Fire: GAS-ODOR/LEAK	2015-12-10 17:40:00	NORRISTOWN	HAWS AVE	1

Basic Questions

What are the top 5 zipcodes for 911 calls?

```
In [6]: df['zip'].value_counts().head(5)

Out [6]:
```

zip	count
19401.0	6979
19464.0	6543
19493.0	4854
19446.0	4748
19406.0	3174

What are the top 5 townships (twp) for 911 calls?

```
In [7]: df['twp'].value_counts().head(5)

Out [7]:
```

twp	count
LOWER MERION	8443
ABINGTON	5977
NORRISTOWN	5890
UPPER MERION	5227
CHELSEA/HAW	4575

Take a look at the 'title' column, how many unique title codes are there?

```
In [8]: df['title'].nunique()

Out [8]: 110
```

Creating new features

In the titles column there are "Reasons/Departments" specified before the title code. These are EMS, Fire, and Traffic. Use .apply() with a custom lambda expression to create a new column called "Reason" that contains this string value.

For example, if the title column value is EMS: BACK PAINS/INJURY , the Reason column value would be EMS.

```
In [9]: df['Reason'] = df['title'].apply(lambda title: title.split(':')[0])

What is the most common Reason for a 911 call based off of this new column?
```

```
In [10]: df['Reason'].value_counts()

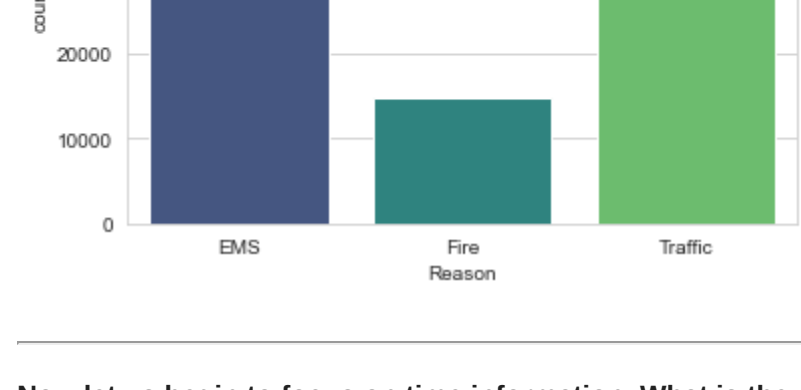
Out [10]:
```

Reason	count
EMS	48877
Traffic	35695
Fire	14929

Now use seaborn to create a countplot of 911 calls by Reason.

```
In [11]: sns.countplot(x='Reason', data=df, palette='viridis')

Out [11]: <matplotlib.axes._subplots.AxesSubplot at 0x1dc0b7e3f78>
```



Now let us begin to focus on time information. What is the data type of the objects in the timeStamp column?

```
In [12]: type(df['timeStamp'].iloc[0])

Out [12]: str

You should have seen that these timestamps are still strings. Use pd.to_datetime to convert the column from strings to DateTime objects.

df['timeStamp'] = pd.to_datetime(df['timeStamp'])

You can now grab specific attributes from a DateTime object by calling them. For example:
```

```
time = df['timeStamp'].iloc[6]
time.hour
```

You can use Jupyter's tab method to explore the various attributes you can call. Now that the timestamp column are actually DateTime objects, use .apply() to create 3 new columns called Hour, Month, and Day of Week. You will create these columns based off of the timeStamp column, reference the solutions if you get stuck on this step.

```
In [14]: df['Hour'] = df['timeStamp'].apply(lambda time: time.hour)
df['Month'] = df['timeStamp'].apply(lambda time: time.month)
df['Day of Week'] = df['timeStamp'].apply(lambda time: time.dayofweek)

Notice how the Day of Week is an integer 0-6. Use the .map() with this dictionary to map the actual string names to the day of the week:
```

```
dmap = {0:'Mon', 1:'Tue', 2:'Wed', 3:'Thu', 4:'Fri', 5:'Sat', 6:'Sun'}

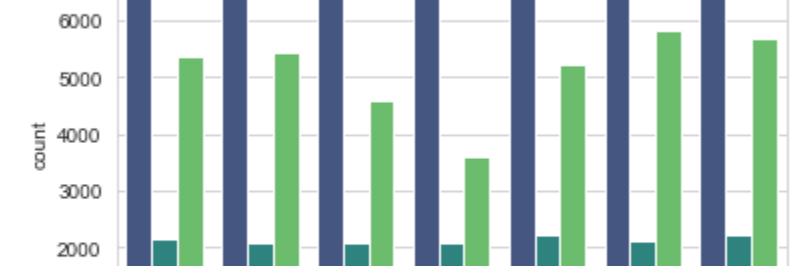
In [15]: dmap = {0:'Mon', 1:'Tue', 2:'Wed', 3:'Thu', 4:'Fri', 5:'Sat', 6:'Sun'}

In [16]: df['Day of Week'] = df['Day of Week'].map(dmap)
```

Now use seaborn to create a countplot of the Day of Week column with the hue based off of the Reason column.

```
In [17]: sns.countplot(x='Day of Week', data=df, hue='Reason', palette='viridis')
# To relocate the legend
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)

Out [17]: <matplotlib.legend.Legend at 0x1dccc0d08b0>
```



Now do the same for Month:

```
In [18]: sns.countplot(x='Month', data=df, hue='Reason', palette='viridis')
# To relocate the legend
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)

Out [18]: <matplotlib.legend.Legend at 0x1dccc4a1af0>
```



Did you notice something strange about the Plot?

```
In [19]: # It is missing some months! 9, 10, and 11 are not there.

You should have noticed it was missing some Months, let's see if we can maybe fill in this information by plotting the information in another way, possibly a simple line plot that fills in the missing months, in order to do this, we'll need to do some work with pandas...

Now create a groupby object called byMonth, where you group the DataFrame by the month column and use the count() method for aggregation. Use the head() method on this returned DataFrame.
```

```
In [20]: byMonth = df.groupby('Month').count()
byMonth.head()


Out [20]:
```

Month	lat	lng	desc	zip	title	timeStamp	twp	addr	e	Reason	Hour	Day of Week
1	13205	13205	13205	11527	13205	13205	13203	13096	13205	13205	13205	13205
2	11467	11467	11467	9930	11467	11467	11465	11396	11467	11467	11467	11467
3	11101	11101	11101	9759	11101	11101	11092	11059	11101	11101	11101	11101
4	11326	11326	11326	9895	11326	11326	11323	11283	11326	11326	11326	11326
5	11423	11423	11423	9949	11423	11423	11420	11378	11423	11423	11423	11423

Now create a simple plot off of the dataframe indicating the count of calls per month.

```
In [21]: # Could be any column
byMonth['twp'].plot()

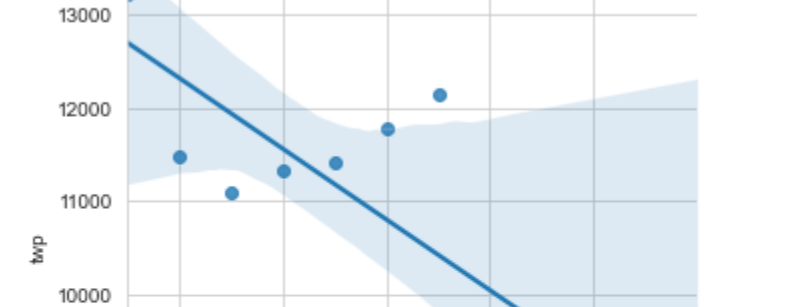
Out [21]: <matplotlib.axes._subplots.AxesSubplot at 0x1dccc15c460>
```



Now see if you can use seaborn's lmplo() to create a linear fit on the number of calls per month. Keep in mind you may need to reset the index to a column.

```
In [22]: sns.lmplo(x='Month', y='twp', data=byMonth.reset_index())

Out [22]: <seaborn.axisgrid.FacetGrid at 0x1dccc13ac0>
```



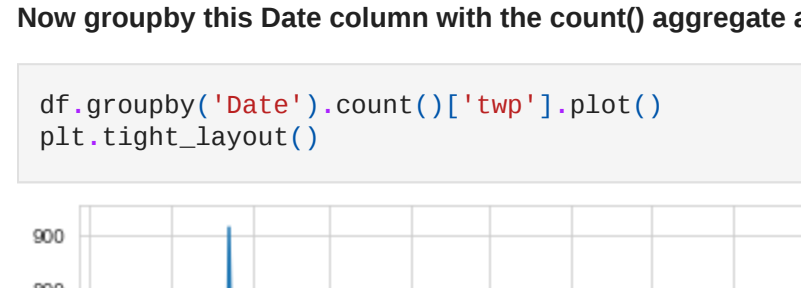
Create a new column called 'Date' that contains the date from the timeStamp column. You'll need to use apply along with the .date() method.

```
In [23]: df['Date'] = df['timeStamp'].apply(lambda t: t.date())

Now groupby this 'Date' column with the count() aggregate and create a plot of counts of 911 calls.

In [24]: df.groupby('Date').count()[['twp']].plot()
plt.tight_layout()

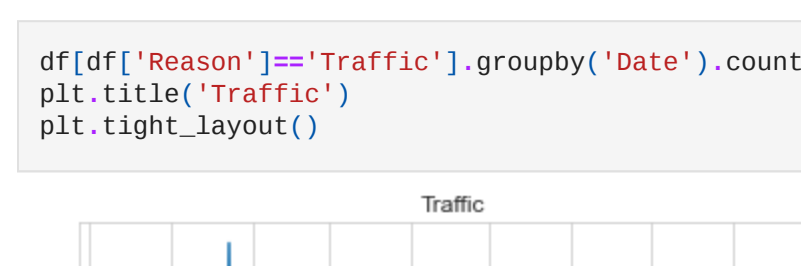
Out [24]:
```



Now recreate this plot but create 3 separate plots with each plot representing a Reason for the 911 call

```
In [25]: df[df['Reason']=='Traffic'].groupby('Date').count()[['twp']].plot()
plt.title('Traffic')
plt.tight_layout()

Out [25]:
```



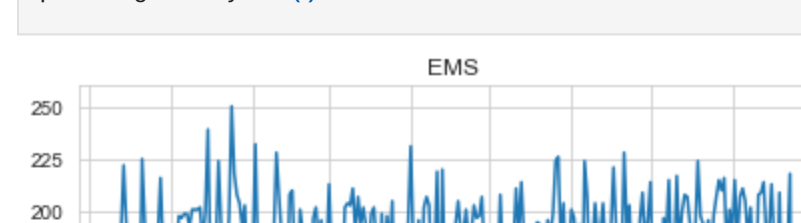
```
In [26]: df[df['Reason']=='Fire'].groupby('Date').count()[['twp']].plot()
plt.title('Fire')
plt.tight_layout()

Out [26]:
```



```
In [27]: df[df['Reason']=='EMS'].groupby('Date').count()[['twp']].plot()
plt.title('EMS')
plt.tight_layout()

Out [27]:
```



Now let's move on to creating heatmaps with seaborn and our data. We'll first need to restructure the dataframe so that the columns become the Hours and the index becomes the Day of the Week. There are lots of ways to do this, but I would recommend trying to combine groupby with an unstack method. Reference the solutions if you get stuck on this!

```
In [28]: dayHour = df.groupby(by=['Day of Week', 'Hour']).count()[['Reason']].unstack()
dayHour.head()

Out [28]:
```


Day of Week	Hour	0	1	2	3	4	5	6	7	8	9	...	14	15	16	17	18	19	20	21	22	23
Fri	275	235	191	175	201	194	372	598	742	752	...	932	980	1039	980	820	696	667	559	514	474	
Mon	282	221	201	194	204	267	397	653	819	786	...	869	913	989	997	885	746	613	497	472	325	
Sat	375	301	263	260	224	231	257	391	459	640	...	789	796	848	757	778	696	628	572	506	467	
Sun	383	306	286	268	242	240	300	402	483	620	...	684	691	663	714	670	655	537	461	415	330	
Thu	278	202	233	159	182	203	362	570	777	828	...	876	969	935	1013	810	698	617	553	424	354	

5 rows x 24 columns

Now create a HeatMap using this new DataFrame.

```
In [29]: plt.figure(figsize=(12,6))
sns.heatmap(dayHour, cmap='viridis')

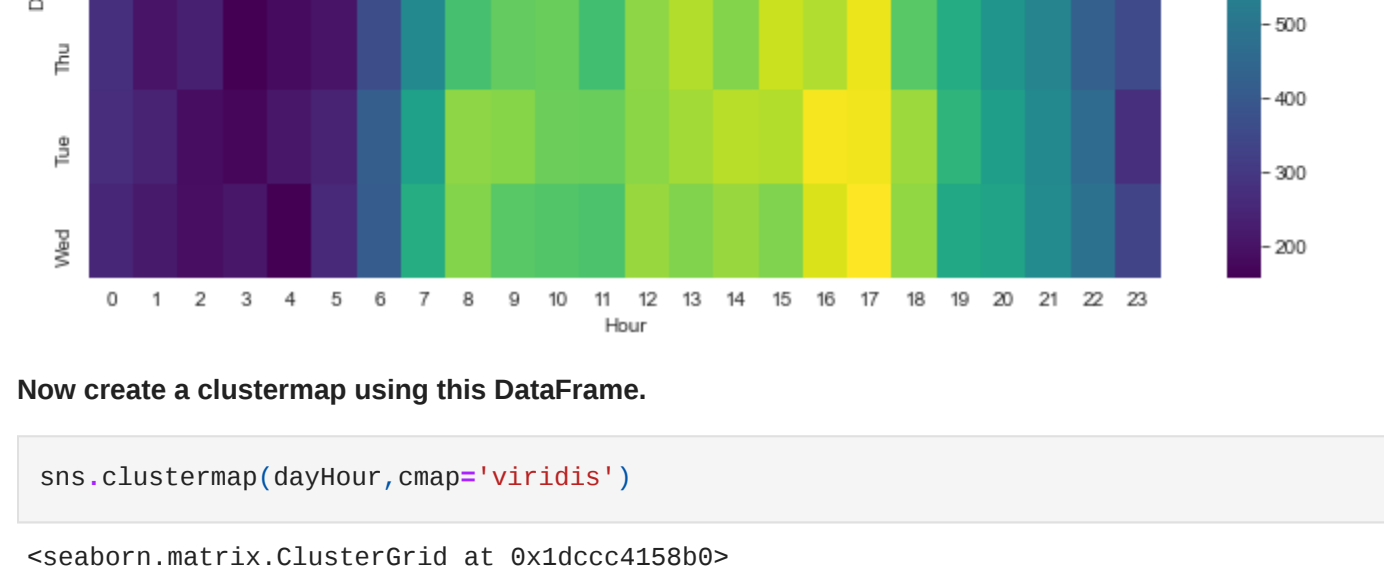
Out [29]: <matplotlib.axes._subplots.AxesSubplot at 0x1dccc41fc0>
```



Now create a clustermap using this DataFrame.

```
In [30]: sns.clustermap(dayHour, cmap='viridis')

Out [30]: <seaborn.matrix.ClusterGrid at 0x1dccc4158b0>
```



Now repeat these same plots and operations, for a DataFrame that shows the Month as the column.

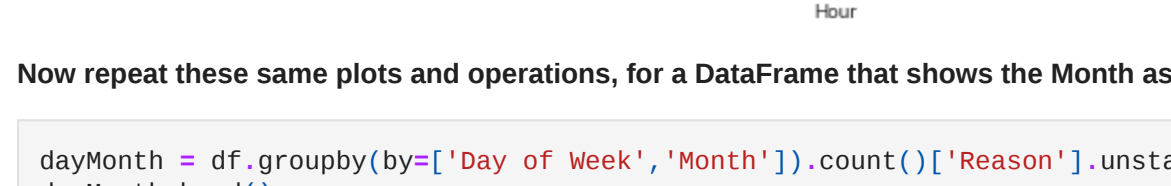
```
In [31]: dayMonth = df.groupby(by=['Day of Week', 'Month']).count()[['Reason']].unstack()
dayMonth.head()

Out [31]:
```

Day of Week	Month	1	2	3	4	5	6	7	8	12
Fri	1970	1581	1525	1598	1730	1649	2045	1310	1065	
Mon	1727	1964	1535	1958	1779	1617	1692	1511	1257	
Sat	2291	1441	1266	1734	1444	1388	1695	1099	978	
Sun	1960	1229	1102	1488	1424	1333	1672	1021	907	
Thu	1584	1596	1900	1601	1590	2065	1646	1230	1266	

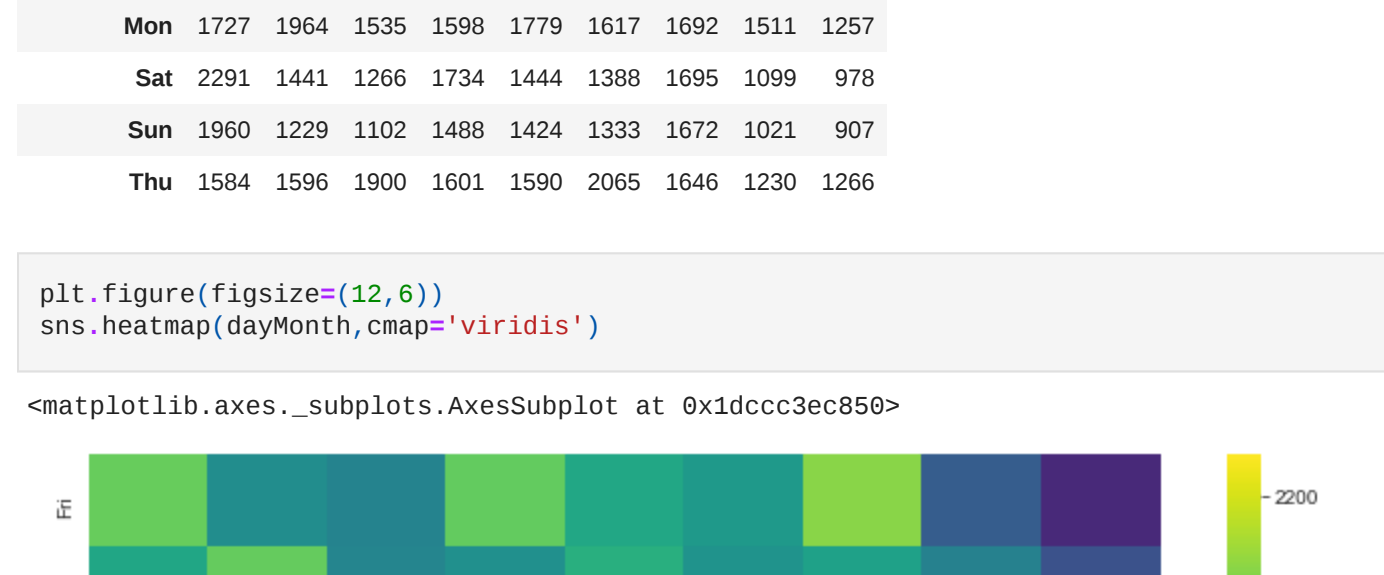
```
In [32]: plt.figure(figsize=(12,6))
sns.heatmap(dayMonth, cmap='viridis')

Out [32]: <matplotlib.axes._subplots.AxesSubplot at 0x1dccc3ec850>
```



```
In [33]: sns.clustermap(dayMonth, cmap='viridis')

Out [33]: <seaborn.matrix.ClusterGrid at 0x1dccc0b9ec0>
```



Continue exploring the Data however you see fit!

Great Job!